Aalto University
School of Electrical
Engineering

Master's Programme in Computer, Communication and Information Sciences

MPC in practice

MPC in practice

Tianxing Wu

Master's Thesis
2024

© 2024

This work is licensed under a Creative Commons "Attribution-NonCommercial-ShareAlike 4.0 International" license.

Aalto University
School of Electrical Engineering

Author Tianxing Wu

Title MPC in practice — MPC in practice

Degree programme Computer, Communication and Information Sciences

Major Algorithms

Supervisor Prof. Jukka Suomela

Advisor Hossein Vahidi

Date 10 July 2024 **Number of pages** 15 **Language** English

Abstract

The purpose of this thesis is to study how congested clique algorithms perform on modern supercomputers.

A representative algorithm $MST_\log\log n$ for Congested Clique is studied. The algorithm is implemented on Puhti, a Finnish supercomputer hosted by CSC. The congested clique setting is simulated with the Message Passing Interface (MPI) and CUDA.

MPI is a message-passing library interface specification that uses a message-passing model for parallel computing.

result: TODO

Keywords Distributed computing, Supercomputing, Congested clique, MST, Puhti, MPI, CUDA, MPC

Preface

I want to thank Professor Jukka Suomela and my advisor Hossein Vahidi for their guidance.

I also want to thank my friends for keeping me sane and alive.

Otaniemi, 10 July 2024

Tianxing Wu

Contents

Abstract 3

Preface 4

Contents 5

Symbols and abbreviations 6

1 Introduction 7

2 Preliminaries 8
- 2.1 Congested clique . . . 8
- 2.2 The algorithm $MST_\log\log n$ . . . 8
- 2.3 The super-computer Puhti . . . 8
- 2.4 MPI . . . 8
- 2.5 CUDA . . . 8

3 Congested clique to MPI 9
- 3.1 Congested Clique simulation . . . 9
- 3.2 Simplification on the first round . . . 9
- 3.3 Union-find . . . 10
- 3.4 Random vertex partitioning . . . 10
- 3.5 Memory limit and graph generation . . . 10

4 MPI Algorithm performance 11
- 4.1 $MST_\log\log n$ time composition . . . 11
- 4.2 performance compared to sequential algorithm . . . 11
- 4.3 Scalability . . . 12
- 4.4 Python vs C++ . . . 12

5 Congested clique to CUDA 13
- 5.1 Congested Clique simulation . . . 13

6 CUDA performance 14

References 15

Symbols and abbreviations

Abbreviations

$MST_\log\log n$ The algorithm described in [1]

1 Introduction

Congested Clique, described in [2] is a mathematical model for distributed computing. It features n inter-connected computers in a clique network, sending messages in a synchronized manner. In each round, each computer sends $O(\log n)$ to each other computer. The computation efficiency is measured by the communication rounds. In this work, the congested clique algorithm that solves the (Minimum spanning tree)MST in $O(\log \log n)$ rounds [1] is studied and implemented on the supercomputer Puhti [3]. We simulated the congested clique model with MPI and a CUDA GPU to run $MST_\log \log n$.

The purpose of this work is to identify the connection between $MST_\log \log n$ and modern distributed computing hardware and tools. The relations allow us to understand the bigger picture as $MST_\log \log n$ is representative of the congested clique model

The representativeness of $MST_\log \log n$ can be shown in the following ways:

- $MST_\log \log n$ exploit the full connectivity of the graph. In each round, there is a step where each computer sends $O(\log n)$ messages to all the other computers.
- To comply with the $O(\log n)$ limit when sending to a unique leader, large data is spilled to other computers before sending.
- The algorithm needs several rounds before its convergence.

We measure how compatible $MST_\log \log n$ is with the supercomputing infrastructure to find the connections. The congested clique is simulated with MPI and CUDA. When implementing the algorithm pseudo-code into runnable, the compatibility is measured by the effort needed for the conversion. The simulation consists of two parts, the local computation, and the message passing. When implementing $MST_\log \log n$, the coding follows best practices and common sense.

Secondly, how the hardware is limiting the performance of the algorithm. Memory, CPU, GPU, network, which one is the bottleneck when the algorithm scales. How the algorithm performs compared to traditional sequential algorithms. By increasing the number of nodes, will it outperform the sequential algorithm with

- Speed comparing to sequential algorithms.
- Number of edges can be processed compared to sequential algorithms.

Scalability:

- Strong scaling and weak scaling (How having more processes speeds up the computation, which part of the hardware will become bottleneck when scaling continuously).
- Time composition for the run when scaling (MPI_time vs local). Does the time complexity hold?

2 Preliminaries

2.1 Congested clique

what is a round? What is each step in a round? The round is synchronous. The communication bandwidth to each neighbor is $O(\log n)$.

2.2 The algorithm $MST_\log\log n$

In each round, the algorithm consists of the following steps: (TODO)

1. Each vertex v computes the minimum weight edge $e(v, F')$ that connects v to a leader l.

2. The numbers start at 1 with each use of the `enumerate` environment.

3. Another entry in the list

The steps are repeated until the MST is found.

2.3 The super-computer Puhti

Used Slurm workload manager [6], to allocate computing nodes. Each node can have processes. The number of total threads is limited, can get 484 * 2 * 20 for the "M" cluster. Total memory available is 192 GiB. Xeon Gold 6230 is the CPU model (ref needed). The nodes are connected with a 100 Gbps HDR100 link, and the topology is a fat tree with a blocking factor of approximately 2:1. [5]

2.4 MPI

2.5 CUDA

3 Congested clique to MPI

3.1 Congested Clique simulation

$MST_\log\log n$ is simulated with OpenMPI and C++ on CPUs. Multiple CPU threads run cooperatively to simulate the congested clique model. Each thread has its dedicated memory space and computing power. Each thread simulates multiple machines in the congested clique model. The local computation of the congested clique is simulated by each thread's local computation while the communication is handled by OpenMPI.

MPI (Message-passing interface) is a message-passing library interface specification that message-passing model for parallel computing [7]. Open MPI is an open-source Message-passing Interface implementation.

One CPU thread can hold one or multiple vertices. The thread number is limited. With each thread simulating more than one vertex, larger graphs can be simulated. To simulate more than one vertex, each thread holds all the edges adjacent to the vertices it is responsible for and run the algorithm for each vertex.

When running $MST_\log\log n$, in the beginning, each rank receives input to run the algorithm. The input includes a list of vertices the rank is responsible for and all the adjacent edges.

step 1 → MPI_Alltoall (comm1)
step 2 → MPI_Alltoall (comm2)
step 3 → MPI_Gather (comm3)
step 4 → Local computation on rank 0
step 5 → MPI_Broadcast (comm4)
step 6 → Local computation on each thread

To ensure the correctness of the algorithm, the result of $MST_\log\log n$ is compared to the result of the prim's algorithm on the same graph.

3.2 Simplification on the first round

We implement the algorithm with the best practices and common sense. The algorithm is modified to align with the implementation goal.

At the beginning of the first round, there is no spanning tree or clusters formed yet. Each vertex itself is a cluster. In the first step, for each vertex v and edge $e \in E_v$, e is sent over the connection e. Because for every two clusters $F \neq F'$, there is only one unique edge connecting them. In step two, each cluster leader (the vertex itself) sorts all the edges it receives, which is all the edges the vertex is adjacent to check the minimum, sort and take only one as μ is 1.

Essentially, all the above steps find each vertex the minimum weight edge adjacent. Simplification is done to the algorithm during the first rounds. In the simplified version, rank 0 gathers the minimum weighed edges adjacent to each vertex (comm0) and skips steps 1-3 during the first round.

The Simplification has a significant impact on the performance of the algorithm. Because the time taken by MPI communication corresponds to the message size.

During the first round, each vertex is a cluster, and the communication needs are the biggest in the whole algorithm.

3.3 Union-find

When merging edges to the MST, each edge is checked if it is connecting two vertices that are already connected. Thus, a union-find is maintained. If two vertices have the same root, it means they are connected. When an edge is added to the MST, two sets containing the vertices are merged.

3.4 Random vertex partitioning

Each thread holds a list of vertices and all the edges adjacent to the vertices. The adjacency matrix represents a complete graph, and the weight on each edge is generated randomly. The adjacency matrix is partitioned and each thread holds a continuous part of the matrix.

3.5 Memory limit and graph generation

Each node can have at most 40 threads and has a memory limit of 192 GiB. With n threads for a node, each thread can have at most $192/n$ GiB memory.

The graph is a complete graph represented by a square adj matrix. Each edge is assigned a random weigdht. There is a naive approach and a distributed approach to generate the graph.

Naive generation: generate the adj matrix on rank 0 and scatter it to all other ranks with MPI_Scatter. Each rank gets a matrix row. Memory eager ($O(n^2)$ memory needed), can't run large enough graph.

Distributed generation: Each rank generates several rows of the adj matrix. And share the result with the destination peer with MPI_Alltoall. It saves the memory for graph storage and allows the running of larger graphs.

4 MPI Algorithm performance

4.1 $MST_\log\log n$ time composition

8 processes, 16G memory each. As the number of vertices goes up, local computation becomes the dominant part in the time composition.

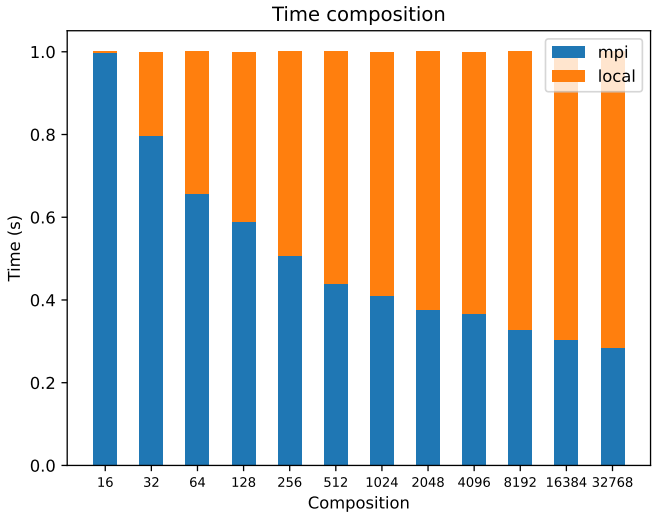

Figure 1: $MST_\log\log n$ time composition implemented by MPI

TODO insert time composition by round chart.

TODO insert graph with 16 processes

The time decreases as the round number goes up. The dominant time in each round is the MPI_Alltoall. Because it has the largest message size.

Each round, MPI time grows linear as the graph size grows. The local computation time has its bottleneck on sorting. So the MPI time grows slower than the local computation.

4.2 performance compared to sequential algorithm

Prim's algorithm is used as the sequential algorithm to compare the performance of $MST_\log\log n$. It is implemented also in C++ and runs on a single thread on Puhti, with 16G memory.

With 16G memory, prim's algorithm can process 32678-vertex complete graph (but not 65356), which has $32678 \times 32677/2 (5.33 \times 10^8)$ edges, using 88 seconds (2). The bottleneck for prim's algorithm to process more edges is to fit the graph to the memory.

Because of the distributed nature of $MST_\log\log n$, larger graphs can be fit into the memory as the number of threads increases. When working on graphs with the same size, the program needs more memory for MPI's buffer compared to the prim's algorithms.

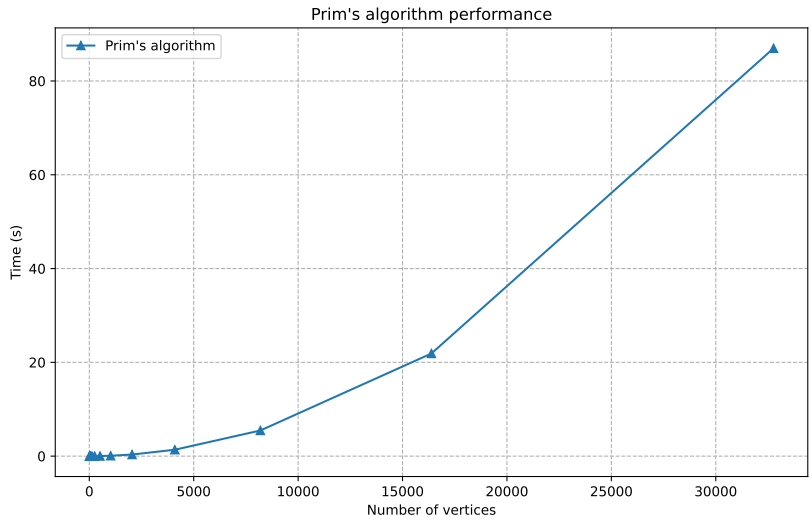

Figure 2: Prim's algorithm performance

As shown in 3, with more than 16 threads, $MST_\log\log n$ can process 65536-vertex complete graph (but not 131072). And $MST_\log\log n$ outperforms the prim's algorithm by speed when there are more than 64 threads.

From 3 as the number of thread increases, the time taken by $MST_\log\log n$ decreases. When

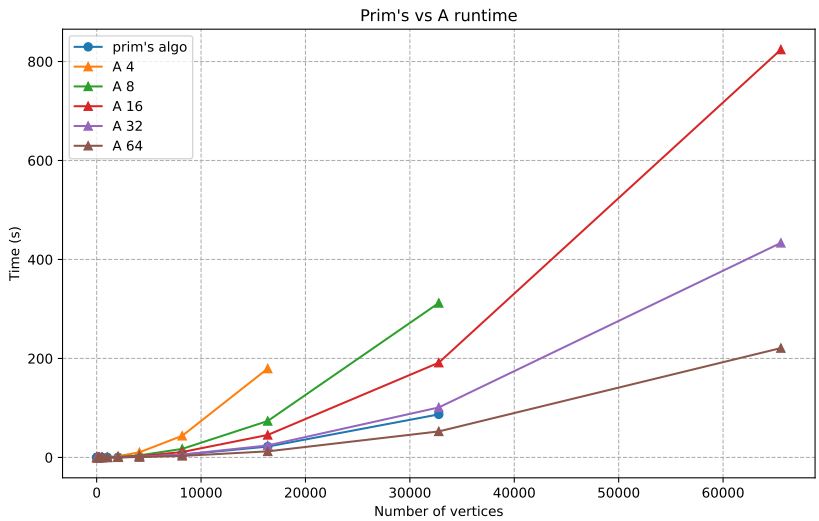

Figure 3: Prim's algorithm compared to $MST_\log\log n$ performance

4.3 Scalability

4.4 Python vs C++

5 Congested clique to CUDA

5.1 Congested Clique simulation

6 CUDA performance

References

This is the list of references to the sources cited in appendix **??**. The list more or less follows the Vancouver style (IEEE). See appendix **??** for a detailed exposition on cross-referencing and bibliography styles. Follow the description there.

[1] Zvi Lotker, Boaz Patt-Shamir, Elan Pavlov, and David Peleg, ”MST Construction in $O(\log \log n)$ Communication Rounds" *SIAM Journal on Computing*, 35(1):120–131, 2005. https://doi.org/10.1137/S0097539704441848

[2] M. Ghaffari, F. Kuhn, and C. Lenzen, "On the power of the congested clique model," *Proceedings of the 2013 ACM Symposium on Principles of Distributed Computing (PODC ’13)*, 2013, pp. 367-376. https://www.researchgate.net/publication/266659337_On_the_power_of_the_congested_clique_model

[3] "Puhti," CSC - IT Center for Science, Accessed: Jul. 16, 2024. [Online]. Available: https://research.csc.fi/-/puhti

[4] "Puhti Service Description," CSC - IT Center for Science, Jun. 2021. [Online]. Available: https://research.csc.fi/documents/48467/640068/Puhti-service description-0621.pdf/e6fd32ae-92ea-2d42-0b6c-9f73c1a4c418?t=1628595451763

[5] "Puhti Computing Systems," CSC - IT Center for Science, Accessed: Jul. 16, 2024. [Online]. Available: https://docs.csc.fi/computing/systems-puhti/

[6] "Slurm Overview," SchedMD, Accessed: Jul. 16, 2024. [Online]. Available: https://slurm.schedmd.com/overview.html

[7] "MPI Forum," Accessed: Jul. 16, 2024. [Online]. Available: https://www.mpi-forum.org/

[8] "Open MPI," Accessed: Jul. 16, 2024. [Online]. Available: https://www.open-mpi.org/